

Homework 3 Solutions

A study was conducted to investigate two treatments for patients suffering from multiple sclerosis. 150 sufferers of the disease were recruited into the study, and 75 were randomized to receive azathioprine (AZ) alone (group 1), and 75 were randomized to receive azathioprine plus methylprednisolone (AZ+MP, group 2). For each participant, a measure of auto-immunity, azathioprine APCR, was planned at clinic visits at baseline (time 0, at initiation of treatment) and at 3, 6, 9, 12, 15, and 18 months thereafter. Multiple sclerosis (MS) affects the immune system. Low values of APCR (approaching 0) are evidence that immunity is improving, which is hopefully associated with a better prognosis for sufferers of MS. Also recorded for each subject was age at entry into the study and an indicator of whether or not the subject had had previous treatment with either of the study agents (0=no, 1=yes). The average age of the men across both treatment groups was 50.45, with SD 6.69. The primary scientific aims of the study are to investigate whether (i) both treatments (AZ or AZ+MP) lower APCR over the 18 month period and (ii) whether treatment with AZ+MP results in different immune system response than does AZ alone and, if so, how it is different in terms of response over time. It was also suspected that a subject's age and prior history might be related to their APCR level at baseline and to the rate at which APCR changes during the 18 month period. The square root of APCR is the response variable of interest (square roots were taken so that the APCR observations better satisfy the assumption of normality).

1. (20 points) Discuss the issues of study design and how they would relate to your decision to use a linear mixed model or a pattern mixture model.

Note: this question was suppose to have “covariance pattern analysis” not “pattern mixture model” in the wording. We haven't discussed pattern mixture model's, so that would not make sense for me to ask. Below, I answer the question how I meant it to be (with covariance pattern model).

This study is a randomized longitudinal trial with repeated APCR measurements on the same subjects over time. That makes the observations within a subject correlated, so a linear mixed model is a natural primary analysis tool. It allows us to model subject-specific trajectories over the 18-month follow-up and use all available measurements, even if some participants miss visits.

A covariance pattern analysis could be used since the data are designed balanced. However, there are 6 measurements, which is a lot for a pattern mixture model. That is, if we did a heterogeneous variance that would be 7 variance parameters. Further, if we did an unstructured correlation that would have 21 correlation parameters. This is a lot for only 150 subjects.

Our options boil down to:

- If data are heterogeneous:
 - Linear mixed model: random intercept and slope (4 covariance parameters)
 - Covariance pattern analysis: heterogeneous compound symmetric (8 covariance parameters).
- If the correlation is non-constant:
 - Linear mixed model: random intercept and slope (4 covariance parameters)
 - Covariance pattern analysis: Unstructured (29 covariance parameters).

As we can see, the linear mixed model gives some more reasonable options in this situation.

2. (80 points) Regardless of your answer to (1), complete an analysis of this data using a linear mixed model. This analysis should yield the best possible answer to the study's scientific aims.

This analysis should include:

- Exploratory data analysis (e.g., figures) and discussion of findings.
- How do we test the hypothesis of interest? Will you include any confounders or effect modifiers?
- Model fit: compare different mean and random effect structures.
- Model diagnostics checks.

Part a

The first thing I'd like to do is to look at a plot of the data by group. I'm also going to treat group as a class variable. In SAS, I will use PROC IMPORT to read the CSV file and PROC SG PANEL / PROC SG PLOT for the plots.

```
proc import datafile="afcr.csv" out=afcr dbms=csv replace;
  guessingrows=max;
run;
```

```
proc format;
  value grpfmt 1 = "AZ" 2 = "AZ+MP";
  value prfmt 0 = "No" 1 = "Yes";
run;
```

```
proc sort data=afcr;
  by group id time;
run;
```

```
/* Raw spaghetti plot by group */
proc sgpanel data=afcr;
  panelby group / columns=2;
```

```


series x=time y=afcr / group=id transparency=0.7;
colaxis label="Time (months)";
rowaxis label="sqrt(AFCR)";
run;


/* Group means on the same plot */
proc means data=afcr noprint nway;
class group time;
var afcr;
output out=means_afcr mean=mean_afcr;
run;

data plot_afcr;
set afcr(in=a) means_afcr(in=b);
run;

proc sgplot data=plot_afcr;
series x=time y=afcr / group=id transparency=0.8;
series x=time y=mean_afcr / group=group lineattrs=(thickness=2);
xaxis label="Time (months)";
yaxis label="sqrt(AFCR)";
run;

```





There does not appear to be much of an effect of the group variable. It appears that a linear time trend may be sufficient for this data. Also, there may be more heterogeneity in the later months of the study so a random slope may be required.

Checking baseline covariates


Now let's look at whether age or prior have an impact on the baseline values of the outcome.

```
proc sgplot data=afcr;
  where time = 0;
  scatter x=age1 y=afcr;
```

```

loess x=age1 y=afcr / lineattrs=(thickness=2);
xaxis label="Age at entry";
yaxis label="sqrt(AFCR)";
run;

```



Clearly there is a relationship between age and AFCR. The relationship appears to be pretty linear.

Now let's check the relationship between prior:

```

proc ttest data=afcr;
  where time = 0;
  class prior;

```

```
var afcr;
run;
```

prior	Method	Mean	95% CL Mean		Std Dev	95% CL Std Dev	
0		12.8064	12.1706	13.4422	2.1654	1.7994	2.7197
1		13.7269	13.1445	14.3093	2.3876	2.0407	2.8778
Diff (1-2)	Pooled	-0.9205	-1.7872	-0.0538	2.2990	2.0333	2.6451
Diff (1-2)	Satterthwaite	-0.9205	-1.7730	-0.0680			

Method	Variances	DF	t Value	Pr > t
Pooled	Equal	112	-2.10	0.0376
Satterthwaite	Unequal	104.8	-2.14	0.0346

Again, a clear relationship here.

Let's see if there are relationships between age, prior, and treatment.

```
proc ttest data=afcr;
  where time = 0;
  class group;
  var age1;
run;
```

The TTEST Procedure

Variable: age1

group	Method	N	Mean	Std Dev	Std Err	Minimum	Maximum
1		55	50.5091	6.9118	0.9320	35.0000	65.0000
2		59	50.6441	6.6508	0.8659	32.0000	67.0000
Diff (1-2)	Pooled		-0.1350	6.7779	1.2704		
Diff (1-2)	Satterthwaite		-0.1350		1.2721		

group	Method	Mean	95% CL Mean		Std Dev	95% CL Std Dev	
1		50.5091	48.6406	52.3776	6.9118	5.8188	8.5143

group	Method	Mean	95% CL Mean		Std Dev	95% CL Std Dev	
2		50.6441	48.9109	52.3773	6.6508	5.6301	8.1270
Diff (1-2)	Pooled	-0.1350	-2.6521	2.3822	6.7779	5.9946	7.7985
Diff (1-2)	Satterthwaite	-0.1350	-2.6559	2.3859			
Method	Variances	DF	t Value	Pr > t			
Pooled	Equal	112	-0.11	0.9156			
Satterthwaite	Unequal	110.68	-0.11	0.9157			

```
proc freq data=afcr;
  where time = 0;
  tables group*prior / chisq;
run;
```

The FREQ Procedure

Frequency Percent Row Pct Col Pct	Table of group by prior			
	group	prior		
		0	1	Total
1		19	36	55
		16.67	31.58	48.25
		34.55	65.45	
		40.43	53.73	
2		28	31	59
		24.56	27.19	51.75
		47.46	52.54	
		59.57	46.27	
Total		47	67	114
		41.23	58.77	100.00

Statistics for Table of group by prior			
Statistic	DF	Value	Prob
Chi-Square	1	1.9586	0.1617
Likelihood Ratio Chi-Square	1	1.9672	0.1607
Continuity Adj. Chi-Square	1	1.4620	0.2266
Mantel-Haenszel Chi-Square	1	1.9414	0.1635
Phi Coefficient		-0.1311	
Contingency Coefficient		0.1300	
Cramer's V		-0.1311	
Fisher's Exact Test			
Cell (1,1) Frequency (F)	19		
Left-sided Pr <= F	0.1132		
Right-sided Pr >= F	0.9444		
Table Probability (P)	0.0576		
Two-sided Pr <= P	0.1857		

No relationships. However, even in a randomized trial, it is reasonable to adjust for baseline covariates such as age and prior treatment when they are strongly associated with the outcome. Randomization ensures that treatment assignment is not confounded, so adjustment is not needed to remove bias. However, including prognostic covariates can improve precision by reducing residual variability and accounting for any chance imbalances between groups. In this setting, since age and prior treatment are clearly related to AFCR, adjusting for them will lead to more efficient estimates of the treatment effects without changing their interpretation.

Part b

The main scientific question is whether the two treatment groups differ in how AFCR changes over time. The primary test is therefore the treatment x time interaction. The model should have time, treatment, and treatment x time.

We do not need to adjust for confounders because the study is randomized. So age and prior treatment are not needed to control confounding of treatment assignment. However, they are reasonable to include as baseline prognostic covariates because they may improve precision and account for chance imbalance.

If we thought there were some effect modifiers, we would put them in using interaction terms. Here, if there is scientific reason to believe age or prior affect the rate of change over time, then they can also be included as effect modifiers through interactions with time, treatment, or treatment x time.

Part c

First, let's check whether a random slope is needed:

```
proc mixed data=afcr;  
  class id group prior;  
  model afcr = age1 prior group time group*time / solution ddfm=bw;  
  random intercept / subject=id;  
run;
```

Fit Statistics	
-2 Res Log Likelihood	3282.3
AIC (Smaller is Better)	3286.3
AICC (Smaller is Better)	3286.3
BIC (Smaller is Better)	3292.2

```
proc mixed data=afcr;  
  class id group prior;  
  model afcr = age1 prior group time group*time / solution ddfm=bw;  
  random intercept time / subject=id type=un;  
run;
```

Fit Statistics

Fit Statistics	
-2 Res Log Likelihood	3281.9
AIC (Smaller is Better)	3289.9
AICC (Smaller is Better)	3289.9
BIC (Smaller is Better)	3301.7

Based on this analysis, a random slope is not required (an LRT was insignificant). So, we'll use the model with just the random intercept.

Before checking some different mean specifications, we'll refit the first model using method = ml.

```
proc mixed data=afcr method=ml;
  class id group prior;
  model afcr = age1 prior group time group*time / solution ddfm=bw;
  random intercept / subject=id;
run;
```

Fit Statistics	
-2 Log Likelihood	3257.1
AIC (Smaller is Better)	3273.1
AICC (Smaller is Better)	3273.3
BIC (Smaller is Better)	3296.6

While the linear trend in AFCD by time was apparent, we'll check a profile analysis just to be complete:

```
proc mixed data=afcr method=ml;
  class id group prior timecat;
  model afcr = age1 prior group timecat group*timecat / solution ddfm=bw;
  random intercept / subject=id;
run;
```

Fit Statistics	
-2 Log Likelihood	3251.9
AIC (Smaller is Better)	3287.9
AICC (Smaller is Better)	3288.8
BIC (Smaller is Better)	3340.9

So, the linear model looks to be the best.

Checking for effect modification (optional)

Now, let's check whether effect modification by age or prior is present. First, I'll check this out for age. I'm only going to consider a three-way age1 x group x time interaction since that is the only one that would substantially change our conclusions.

```
proc mixed data=afcr method=ml;
  class id group prior;
  model afcr = age1 prior group time group*time
              age1*time age1*group age1*group*time / solution ddfm=bw;
  random intercept / subject=id;
run;
```

Fit Statistics	
-2 Log Likelihood	3255.4
AIC (Smaller is Better)	3277.4
AICC (Smaller is Better)	3277.7
BIC (Smaller is Better)	3309.7

No effect modification for age. Now, let's check for prior:

```
proc mixed data=afcr method=ml;
  class id group prior;
  model afcr = age1 prior group time group*time
              prior*time prior*group prior*group*time / solution ddfm=bw;
  random intercept / subject=id;
run;
```

Fit Statistics	
----------------	--


Fit Statistics	
-2 Log Likelihood	3252.7
AIC (Smaller is Better)	3274.7
AICC (Smaller is Better)	3275.0
BIC (Smaller is Better)	3307.1

No effect modification for prior. However, it is close and there is a significant interaction. So, if that was your final model then that is OK. Otherwise, we'll go back to the first model we fitted.

Here's the final model:

```
proc mixed data=afcr method=reml;  
  class id group prior;  
  model afcr = age1 prior group time group*time / solution vciry residual  
influence(effect=id est);  
  random intercept / subject=id;  
run;
```


Here's the residual diagnostics:




The residuals look great.

Here's some influential plots:

Restricted Likelihood Distance





ID 11 and 38 stick out some. Both are 35, which is on the younger side in this study and probably giving them more influence (particularly because their influence seems to be on the intercept). I looked over their data, and it looks fine. We should leave them in the data.

Part d

Here's the coefficients for the final model:

Solution for Fixed Effects							
Effect	group	prior	Estimate	Standard Error	DF	t Value	Pr > t
Intercept			5.5772	0.7907	136	7.05	<.0001
age1			0.1620	0.01491	648	10.87	<.0001
prior		0	-0.9569	0.2061	648	-4.64	<.0001
prior		1	0	.	.	.	.
group	1		-0.3266	0.2772	648	-1.18	0.2392
group	2		0	.	.	.	.
time			-0.2473	0.01489	648	-16.60	<.0001
time*group	1		0.08645	0.02112	648	4.09	<.0001
time*group	2		0	.	.	.	.

Type 3 Tests of Fixed Effects				
Effect	Num DF	Den DF	F Value	Pr > F
age1	1	648	118.05	<.0001
prior	1	648	21.56	<.0001
group	1	648	1.39	0.2392
time	1	648	373.78	<.0001
time*group	1	648	16.76	<.0001

Despite the first figure appearing that there was not an impact of the drug, the final model does show that the treatment was effective. Specifically, the AZ + MP group had their azathioprine decrease more rapidly than the group that had AZ alone. In the AZ + MP group we expect that the outcome ($\sqrt{\text{AFCR}}$) will decrease 0.2473 units for every month increase. In the AZ group, the outcome will decrease 0.089 units less than in the AZ + MP group.